

Assignment 2 Report

Subject: ECE 284 - Parallel Computing in Bioinformatics

Submitted by: Ashwin Agesh Somanathan (A59025915)


Parallelizing seed table construction

Q 5)

Part a)

KmerSize = 12

NumBlocks	BlockSize	NumBlocks x BlockSize	kmerPosConcat (ms)	kmerOffsetFill (ms)	kmerPosMask (ms)	Total kernal time (ms)
64	32	64x32	1.3501	46.714	3.4651	51.5292
64	128	64x128	1.4926	44.195	2.0878	47.7754
64	512	64x512	8.5314	48.834	8.5904	65.9558
256	32	256x32	1.4917	15.807	2.1129	19.4116
256	128	256x128	7.4324	17.952	8.2742	33.6586
256	512	256x512	15.284	23.861	23.256	62.401
1024	32	1024x32	7.933	8.727	8.7876	25.4476
1024	128	1024x128	14.318	22.071	24.05	60.439
1024	512	1024x512	35.563	51.592	48.796	135.951



Part b)

Best overall performance seen for the following configuration:

Numblocks (gridsize) = 256

Block Size = 32

Total time = 19.4116ms

Part c)

Relative speedup for the best performing configuration w.r.t. the original sequential case:

Kernel	kmerPosConcat (us)	kmerOffsetFill (us)	kmerPosMask (us)	Total kernal time (us)
Original code (Sequential)	768.57	1661.71	1007.75	3438.03
Fastest parallel	1.4917	15.807	2.1129	19.4116
Speedup	515.23	105.12	476.95	177.11



Part d)

For the configuration 256 x 32, GPU overheads:

Function	Execution time (Avg) (ms)
cudaMalloc	29.822
[CUDA memcpy HtoD]	0.62290
[CUDA memcpy DtoH]	0.0017920
cudaMemcpy	0.290
cudaFree	2.2469

The average cudaMalloc time is of the same order as the total execution time of the 3 kernels in the fastest parallel implementation. The minimum (0.2ms), however, is small. Since memory allocations are done before calling seedTableOnGpu, the runtime of this function is not impacted (26ms timed by timer object).

Part e)



Bonus Questions

A)

The number of entries in the kmerOffset table is a function of kmerSize related as:

$$\text{Number of entries} = 4^{\text{kmerSize}}$$

The base 4 is a consequence of there being 4 bases: A, C, G, T. These can appear in any permutation.

It is possible to find occurrences larger than kmerSize for a table.

This can be done by breaking the larger kmer into smaller parts and finding the kmer entry that starts or ends with the sub-part. Then for the two kmer entries, find the entries in the position table where they sit in consecutive positions, signifying they are placed one after the other.

Of course, the order of the sub-parts needs to be taken into account.

For example:

If we need to find a kmer AGTC in a table of kmerSize 2, we can break the search item into 2:

AG and TC

Then look in the position table where all they occur. If there are cases where TC is right after AG, then AGTC has been found.

B)

For large kmer sizes, the method mentioned in part (A) can be employed. Build a table for a smaller kmerSize, say 10 and break the actual kmer into parts.